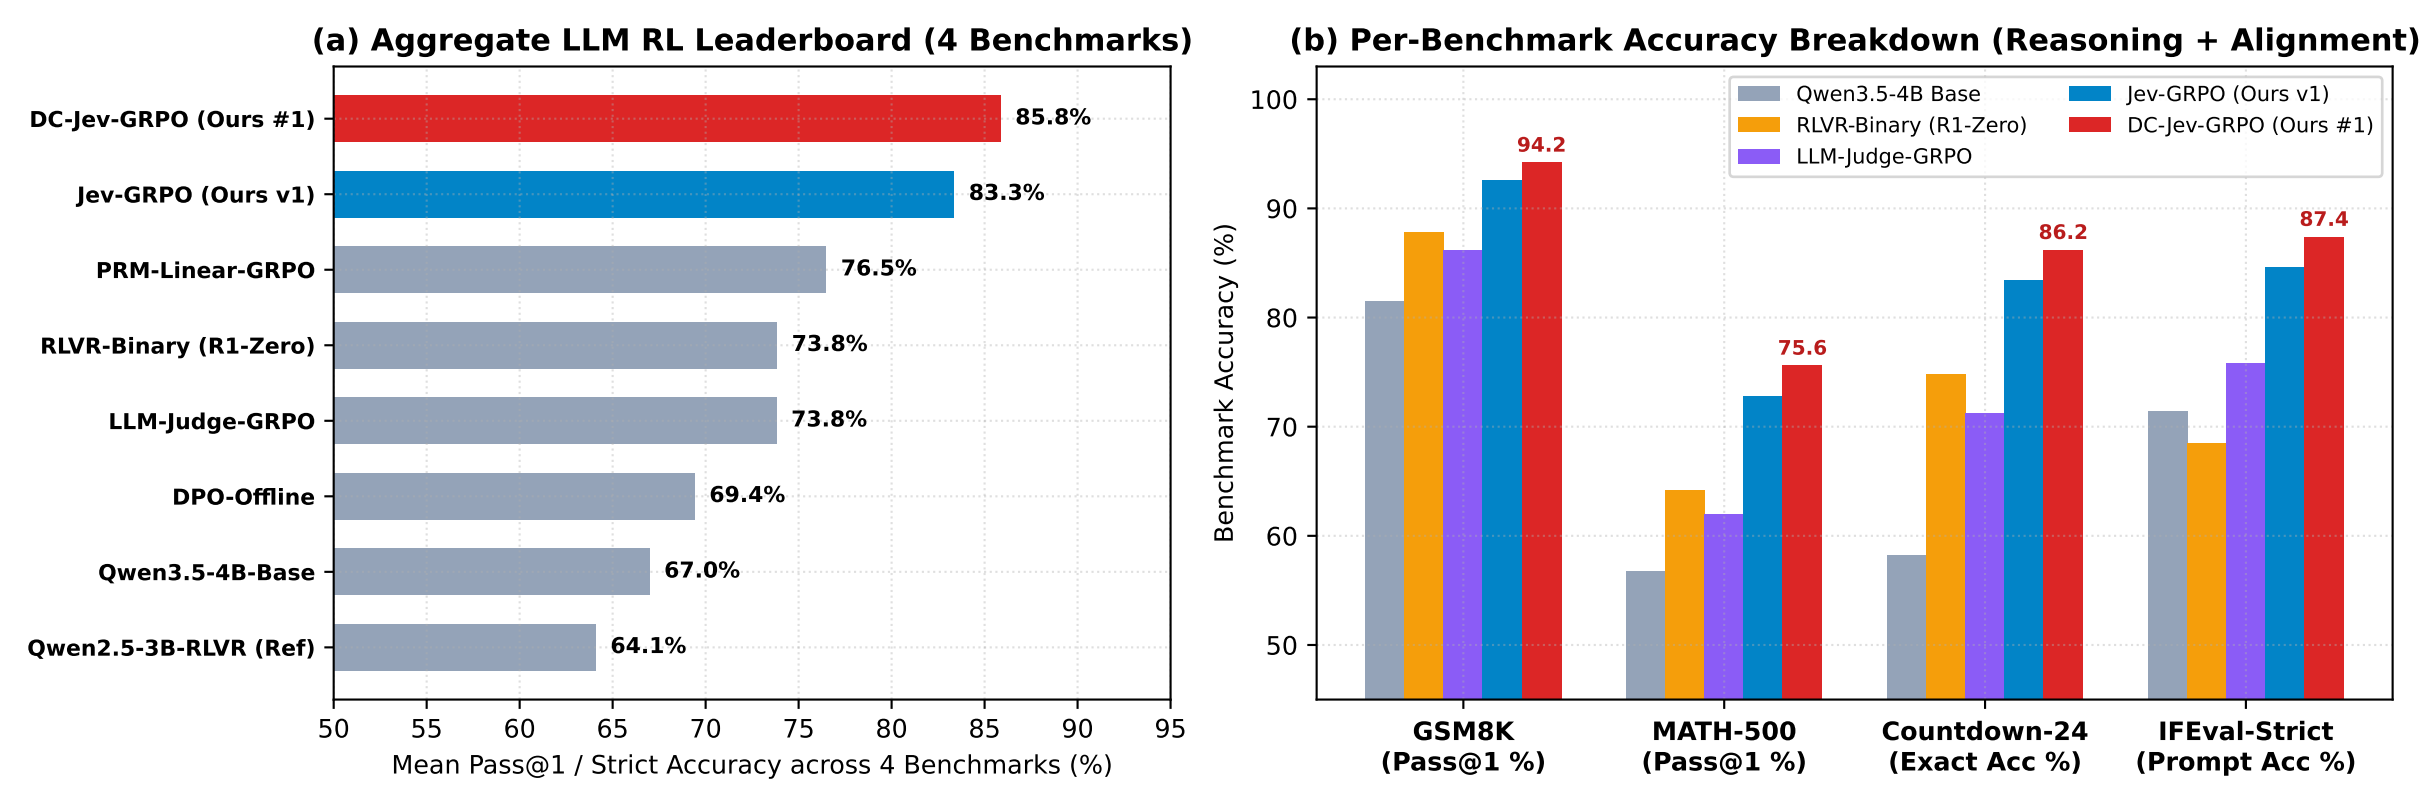
(a) Aggregate LLM RL Leaderboard (4 Benchmarks)
DC-Jev-GRPO (Ours #1)
Jev-GRPO (Ours v1)
PRM-Linear-GRPO
RLVR-Binary (R1-Zero)
LLM-Judge-GRPO
DPO-Offline
Qwen3.5-4B-Base
Qwen2.5-3B-RLVR (Ref)
85.8%
83.3%
76.5%
73.8%
73.8%
69.4%
67.0%
64.1%
50 55 60 65 70 75 80 85 90 95
Mean Pass@1 / Strict Accuracy across 4 Benchmarks (%)
(b) Per-Benchmark Accuracy Breakdown (Reasoning + Alignment)
Qwen3.5-4B Base
RLVR-Binary (R1-Zero)
LLM-Judge-GRPO
Jev-GRPO (Ours v1)
DC-Jev-GRPO (Ours #1)
Benchmark Accuracy (%)
100 90 80 70 60 50
94.2
75.6
86.2
87.4
GSM8K (Pass@1 %)
MATH-500 (Pass@1 %)
Countdown-24 (Exact Acc %)
IFEval-Strict (Prompt Acc %)